A FAIR and project-oriented template for open science data workflows

Kauê de Sousa[a,b,*], Marie-Angélique Laporte[c]

[a] *Digital Inclusion, Bioversity International, Montpellier, France*
[b] *Department of Agricultural Sciences, University of Inland Norway, Hamar, Norway*
[c] *Performance Innovation and Strategic Analysis for Impact (PISA4), Bioversity International, Montpellier, France*

Abstract

Managing research data in accordance with the FAIR principles (Findable, Accessible, Interoperable, and Reusable) is essential for transparency, reproducibility, and effective scientific collaboration (Wilkinson et al., 2016). Reproducibility, the ability to repeat an analysis, and replicability, the ability to repeat an experiment under similar conditions (Stevens, 2017), are foundational to trustworthy and collaborative science (Powers and Hampton, 2019; Munafò et al., 2017). These practices let researchers revisit analyses after long periods and let peers verify results or generate new insights from existing data. Researchers now also work with AI assistants, which can only reuse an analysis well when its structure, reasoning and sources are written down explicitly. We present a structured and reusable repository template that supports FAIR data workflows, promotes reproducibility and replicability, and makes each analysis traceable by people and AI assistants alike. Version 1.3 adds methodological decision records, provenance links from outputs to data, protocols and evidence, automated tests, pinned package versions, and instructions for AI assistants.

Keywords: data management, metadata, FAIR principles, reproducibility, open science, AI-ready documentation

Suggested citation

> de Sousa, K.; Laporte, M.-A. (2026) A FAIR and project-oriented template for open science data workflows. Version 1.3. Zenodo. `https://doi.org/10.5281/zenodo.15224389`.

Version 1.2 of this paper is available at `https://hdl.handle.net/10568/174310`.

*Corresponding author
Email addresses: `k.desousa@cgiar.org` (Kauê de Sousa), `m.a.laporte@cgiar.org` (Marie-Angélique Laporte)

Statement of need

Data-driven research faces several challenges, including inadequate data documentation, inconsistent structure, and poor reproducibility (Wilkinson et al., 2016; Munafò et al., 2017). Researchers often struggle with these aspects for lack of practical examples and accessible tools. To address these issues, we created a project-oriented, easy-to-use template for data analysis. It streamlines metadata management, facilitates reproducibility, and aligns projects with the FAIR principles.

A second need has grown since the first release. Researchers increasingly ask AI assistants to explain, rerun or adapt analyses. An assistant can reuse code, but it cannot recover the reasoning behind it, the conditions under which the data were collected, or the limits of what the results show, unless these are documented. Without that documentation, an assistant may present a method as validated because it appears in a repository, or invent metadata it cannot find. Version 1.3 addresses this need by recording why each analysis was done, linking every output to its sources, and giving assistants explicit rules.

Repository structure

The template repository includes clear, intuitive folder structures and essential resources.

```
template-repo-data-analysis/
|-- data/                          # Anonymized raw and cleaned datasets
|-- decisions/                     # One decision record per script
|    |-- _template.md              # Template for a new record
|-- docs/                          # Reports, papers or other documentation
|-- metadata/                      # Metadata files and templates
|    |-- project-metadata.xlsx     # Project metadata, edited by hand
|    |-- project-metadata.json     # DataCite metadata, generated
|    |-- provenance.json           # Links outputs to code, data and evidence
|    |-- example-metadata-*.csv    # Data dictionary of the example dataset
|-- output/                        # Model results, figures, tables
|-- script/                        # Analysis scripts, run in numeric order
|    |-- helper/                   # Environment setup, metadata generation
|-- tests/                         # One test per analysis script
|-- .github/workflows/check.yml    # Runs all tests on GitHub
|-- AGENTS.md                      # Short instructions for AI agents
|-- SKILL.md                       # Skill for AI assistants
|-- llms.txt                       # Entry point for AI tools on the web
|-- CITATION.cff                   # Citation metadata for GitHub and Zenodo
|-- LICENSE                        # Terms for reusing the work
|-- template-repo-data-analysis.Rproj
|-- README.md                      # Project overview
```

Key features

- **FAIR-aligned metadata.** Metadata files (.xlsx, .json) follow the Zenodo (`https://zenodo.org`) and DataCite (Robinson-Garcia et al., 2017) schemas. Data dictionaries describe every variable and use Crop Ontology terms where they exist.
- **Automated metadata workflows.** R scripts (R Core Team, 2024) validate the metadata and convert them to `CITATION.cff` and DataCite JSON.
- **Decision records.** Each analysis script has a record stating the scientific question, the method and why it was chosen, its assumptions, alternatives and diagnostics, and what the results do and do not support.

- **Provenance.** A machine-readable file links each output to the script that produced it, the input data, the protocol behind the data, and the published evidence.
- **Reproducibility.** Package versions are pinned with renv (`https://rstudio.github.io/renv/`), each script records the software versions of its last run, and each script has a test with documented expected results.
- **Continuous checks.** GitHub Actions runs every test on each push, so the repository always shows whether its computations still reproduce.
- **AI readiness.** `SKILL.md`, `AGENTS.md` and `llms.txt` tell AI assistants how the project is organized, which conventions to follow, and which rules to respect.
- **Citation integration.** A `CITATION.cff` file lets GitHub and Zenodo generate citations and DOIs.
- **Open and collaborative.** Open licensing (CC BY 4.0, `https://creativecommons.org`) promotes collaboration, reuse, and attribution.

AI-ready documentation

The template belongs to ai-ready-documentation (`https://github.com/AgrDataSci/ai-ready-documentation`), a family of linked repositories that make the tricot approach to on-farm testing usable by people and AI assistants. Each repository in the family covers one layer, namely the published evidence, the field protocols, the trial data and the analytical approaches, and each shares the same skeleton of `SKILL.md`, `README.md`, `CITATION.cff` and `LICENSE`. The template is not itself a layer of evidence. It is the scaffold from which any layer, or any research project, can start.

Two principles guide the design. First, reproducing a computation is different from validating an interpretation. A passing test shows that the code reruns and gives the expected numbers, while the decision record states whether the method suits the question and what the results cannot support. Second, missing information is stated rather than filled in. The provenance file uses `null` for unknown values, and the instructions forbid an assistant from guessing a DOI, ORCID, license, funder or date.

Worked example

The template ships with a complete example that shows every component together. The data are rankings from 150 farmers in Uganda who each watched three of eight video-based maize hybrid concepts and named the most and least preferred, drawn from a product concept testing study in Kenya and Uganda (Donovan et al., 2025). A data dictionary describes each variable. The analysis script fits a Plackett-Luce model to the rankings (Turner et al., 2020) and saves the estimates and a plot. Its decision record explains why this model fits tricot rankings, states its assumptions and alternatives, and limits the conclusions to relative preference within the sample, not adoption or market size.

The test checks that the 150 rankings are complete, that all eight concepts receive finite estimates, and that the most and least preferred concepts match an independent fit of the same data. Drought Avoidance, Intercropping and Resilient are the most preferred concepts and Feed Variety the least preferred, consistent with the full study (Donovan et al., 2025). The order within the top three is close, so the test checks the group rather than the exact order.

Implementation and usage

Researchers adopt the template in these steps.

1. Copy the repository from GitHub (`https://github.com/AgrDataSci/template-repo-data-analysis`) and pin package versions with the setup helper script.
2. Add the data and describe each variable in a data dictionary.
3. Write the analysis in numbered scripts, each with a decision record and a test.
4. Register inputs, outputs, protocols and evidence in the provenance file.

5. Fill in the provided `project-metadata.xlsx` template and run the R scripts for validation and metadata generation.
6. Publish through Zenodo for long-term archiving and DOI assignment.

What is new in version 1.3

Version 1.3 adds decision records, the provenance file, tests with continuous checks on GitHub, pinned package versions, and instructions for AI assistants. The example script now checks the rankings before fitting the model and saves its outputs. The documentation was revised throughout, and the template now links to the other repositories of ai-ready-documentation.

Conclusion

Structured, FAIR-aligned workflows enhance reproducibility, collaboration, and transparency in research. By recording why each analysis was done and where every result comes from, this template also lets AI assistants reuse the work without overstating it. It provides a practical starting point that promotes these practices from the outset of a research project.

Acknowledgments

This work has been supported by the projects 1000FARMS (INV-031561) funded by the Gates Foundation, and BOLD Work Package 7 (Crop Trust Ref: CONT-1522) funded by the Norwegian Agency for Development Cooperation (Norad) through The Global Crop Diversity Trust.

References

J. Donovan, P. Rutsaert, H. Mawia, K. de Sousa, and J. van Etten. Farmers' preferences for the next generation of maize hybrids. application of product concept testing in kenya and uganda. *Experimental Agriculture*, 61:e5, 2025. doi: 10.1017/S001447972500002X. URL `https://doi.org/10.1017/S001447972500002X`.

Marcus R Munafò, Brian A Nosek, Dorothy V M Bishop, Katherine S Button, Christopher D Chambers, Nathalie Percie du Sert, Uri Simonsohn, Eric-Jan Wagenmakers, Jennifer J Ware, and John P A Ioannidis. A manifesto for reproducible science. *Nature Human Behaviour*, 1(1):21, 2017. ISSN 2397-3374. doi: 10.1038/s41562-016-0021. URL `https://doi.org/10.1038/s41562-016-0021`.

Stephen M Powers and Stephanie E Hampton. Open science, reproducibility, and transparency in ecology. *Ecological Applications*, 29(1):e01822, jan 2019. ISSN 1051-0761. doi: 10.1002/eap.1822. URL `https://doi.org/10.1002/eap.1822`.

R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2024. URL `https://www.R-project.org/`.

Nicolas Robinson-Garcia, Philippe Mongeon, Wei Jeng, and Rodrigo Costas. Datacite as a novel bibliometric source: Coverage, strengths and limitations. *Journal of Informetrics*, 11(3):841–854, August 2017. ISSN 1751-1577. doi: 10.1016/j.joi.2017.07.003. URL `http://dx.doi.org/10.1016/j.joi.2017.07.003`.

Jeffrey R. Stevens. Replicability and Reproducibility in Comparative Psychology. *Frontiers in Psychology*, 8:862, may 2017. ISSN 1664-1078. doi: 10.3389/fpsyg.2017.00862. URL `http://journal.frontiersin.org/article/10.3389/fpsyg.2017.00862/full`.

Heather L. Turner, Jacob van Etten, David Firth, and Ioannis Kosmidis. Modelling rankings in R: the PlackettLuce package. *Computational Statistics*, 2020. doi: 10.1007/s00180-020-00959-3. URL `https://doi.org/10.1007/s00180-020-00959-3`.

Mark D. Wilkinson, Michel Dumontier, IJsbrand Jan Aalbersberg, Gabrielle Appleton, Myles Axton, Arie Baak, Niklas Blomberg, Jan-Willem Boiten, Luiz Bonino da Silva Santos, Philip E. Bourne, Jildau Bouwman, Anthony J. Brookes, Tim Clark, Mercè Crosas, Ingrid Dillo, Olivier Dumon, Scott Edmunds, Chris T. Evelo, Richard Finkers, Alejandra Gonzalez-Beltran, Alasdair J.G. Gray, Paul Groth, Carole Goble, Jeffrey S. Grethe, Jaap Heringa, Peter A.C 't Hoen, Rob Hooft, Tobias Kuhn, Ruben Kok, Joost Kok, Scott J. Lusher, Maryann E. Martone, Albert Mons, Abel L. Packer, Bengt Persson, Philippe Rocca-Serra, Marco Roos, Rene van Schaik, Susanna-Assunta Sansone, Erik Schultes, Thierry Sengstag, Ted Slater, George Strawn, Morris A. Swertz, Mark Thompson, Johan van der Lei, Erik van Mulligen, Jan Velterop, Andra Waagmeester, Peter Wittenburg, Katherine Wolstencroft, Jun Zhao, and Barend Mons. The fair guiding principles for scientific data management and stewardship. *Scientific Data*, 3(1), March 2016. ISSN 2052-4463. doi: 10.1038/sdata.2016.18. URL `http://dx.doi.org/10.1038/sdata.2016.18`.